

I've checked your model compilation and callbacks, and they seem to be generally well-configured. Here's a breakdown with some minor suggestions:

### Callbacks:

- **ModelCheckpoint:** Good. You're saving the best model based on validation accuracy.
- **EarlyStopping:** Good. You're using early stopping to prevent overfitting and save time.
- **ReduceLROnPlateau:** Good. You're reducing the learning rate when validation accuracy plateaus, which can help the model converge better.
- **CSVLogger:** Good, you are logging the results.

### Model Compilation:

- **loss:** You're using 'categorical\_crossentropy', which is correct for multi-class classification with one-hot encoded labels (which you have with `label_mode='categorical'` in your data loading).
- **loss\_weights:** You've set `loss_weights=0.3`. This is unusual for a single loss function. `loss_weights` is typically used when you have *multiple* loss functions (e.g., from multiple outputs in a multi-task learning scenario). If you only have one loss function, you should remove the `loss_weights` argument, or set it to 1.
- **optimizer:** You're using SGD with momentum and Nesterov acceleration, which is a good choice.
- **metrics:** You're tracking 'accuracy', which is appropriate for your classification task.

### Suggestions:

- **Remove loss\_weights:** Unless you have a specific reason to weight the categorical cross-entropy loss by 0.3, remove the `loss_weights` argument. It's not doing what you want it to do in this case.
- **Learning Rate:** Double check the learning rate. It is not defined in the code you provided. Make sure it is properly defined.

### Corrected Compilation:

Here's how I recommend you compile your model:

```
model.compile(loss='categorical_crossentropy',
              optimizer=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9,
nesterov=True), # set learning rate
              metrics=['accuracy'])
```

### Reasoning for Changes:

- **loss\_weights Removed:** I removed the `loss_weights` argument because you only have one loss function.
- **learning\_rate:** Add learning rate.

Overall, your code is well-structured. Just remove the unnecessary `loss_weights` argument and ensure your learning rate is correctly defined.